

RTI Exercise 1 Analysis

Preliminary Analysis

My first instinct upon examining this dataset was to look at the distributions of all of the attributes independently. Since I'm a visual person, I conducted this analysis by creating bar charts. After that I used association rule mining¹ on a subset of the attributes (age, workclass, education level, married status, occupation, race, sex, hours/week, country and 50k) to identify combinations of attributes that implied other attributes (with confidence over 90% and support over 20%). Some observations follow:

1. This dataset contains considerable skew in a few attributes towards specific values, which may be representative of the US population. For example,
 - Workclass is skewed towards private (69%)
 - Race is skewed towards white (86%)
 - Sex is skewed towards male (67%)
 - Country is overwhelmingly US (90%)
 - Most importantly, over 50k is skewed towards under (76%)

These skews are not necessarily an indication that the data gathering metric is faulty, but are definitely something we should watch out for in our analysis, since over-sampling a population could produce false correlations in the data or mislead our classifier, if the test set does not contain similar skews. Especially for the label attribute (over 50k), an unbalanced training set will cause problems for some classifier algorithms.

2. There is a significant amount of missing data for the occupation attribute, but not nearly as much for any other attribute.
3. The relationship attribute puzzles me, it is not described in the readme and has some odd trends, for instance, there are more than 8 times as many husbands as wives. If this is census data, that probably means that wives fill out the form for their household far more often than husbands, or this dataset over-samples gay men in marriages heavily.
4. Notable significant association rules include
 - (a) over 50k implies white 90% of the time
 - (b) private sector and female implies under 50k 91% of the time
 - (c) never married implies under 50k 95% of the time

¹Using the A priori algorithm, implementation at <https://github.com/asaini/Apriori>

Classifier Results

I decided to clean some features and remove others to decrease noise before giving it to a classifier. Specifically, I converted the capital gain and capital loss values to a nominal attribute: whether they were zero or not, I converted the country attribute to US or not and removed the education number attribute, since it correlated with the education level attribute. Since the balance of positive to negative examples was so skewed, I elected to experiment with several different ratios of positive to negative examples in the training and test sets to see which resulted in the highest prediction accuracy. I also experimented with three different potential models, random forest, naive bayes and logistic regression classifiers.²

The best model and set of metaparameters was a logistic regression classifier where 90% of the positive examples and 56% of the negative examples were in the training set, which gave the training set a 2-1 bias in favor of negative examples. The model's prediction accuracy was 85.2%, the confusion matrix is shown below.

	predicted under 50k	predicted over 50k
actual under 50k	20383	2701
actual over 50k	1423	3299

Table 1: Confusion Matrix for logistic regression classifier

Rather than include the most informative chart I could find, I decided to include what I found to be the most interesting. Figure 1 shows marital status vs. sex and raises a number of questions. There are over ten times as many husbands as wives, they can't all be married to each other, so where are all of the wives? Why are there four times as many unmarried women as men, despite the dataset being majority male? Are the two people who reported that they are male wives transgender people who are reporting their sex assigned at birth or just errors in the form?

²Implementations from scikit-learn <http://scikit-learn.org/stable/>




Figure 1: Intersections of relationship and sex attributes